

# Linear Regression from Scratch: California Housing

Roll Number: g25ait1010

## 1. Assumptions

- Features are standardized to mean 0 and variance 1; target is unscaled.
- Dataset has no missing values and features follow a linear relationship with target.

## 2. Resources Used

- Scikit-learn for data loading and train/test split.
- NumPy for numerical operations.
- Matplotlib for plotting.
- GeeksforGeeks Gradient Descent formulas[1].
- FutureAI blog on MSE and R<sup>2</sup> definitions[2].

## 3. Implementation

```
class LinearRegression:  
    def __init__(self, learning_rate=0.01, n_iterations=1000, lambda_param=0):  
        self.learning_rate = learning_rate  
        self.n_iterations = n_iterations  
        self.lambda_param = lambda_param  
        self.weights = None  
        self.bias = None  
        self.cost_history = []  
  
    def fit(self, X, y):  
        n_samples, n_features = X.shape  
        self.weights = np.zeros(n_features)  
        self.bias = 0  
        for i in range(self.n_iterations):  
            y_pred = X.dot(self.weights) + self.bias  
            cost = np.mean((y_pred - y)**2)  
            dw = (1/n_samples)*(X.T.dot(y_pred-y) + self.lambda_param*self.weights)  
            db = np.mean(y_pred-y)  
            self.weights -= self.learning_rate*dw  
            self.bias -= self.learning_rate*db  
            self.cost_history.append(cost)  
  
    def predict(self, X):  
        return X.dot(self.weights) + self.bias
```

## 4. Evaluation Metrics

Metric	Definition
MSE	Mean Squared Error: average squared error[3]
R <sup>2</sup>	Variance explained: squared correlation[4]

## 5. Experimental Results

### 5.1 Base Model ( $\lambda=0$ )

Dataset	MSE	R <sup>2</sup>
Train	0.5457	0.5765
Test	0.5238	0.6405

### 5.2 Ridge Model ( $\lambda=10$ )

Dataset	MSE	R <sup>2</sup>
Test	0.5179	0.6471

## 6. Output Plots

### 6.1 Learning Curve (MSE vs. Iterations)

Plots cost reduction during training iterations, showing steep drop and plateau.

### 6.2 Actual vs. Predicted

Scatter of predictions vs. actual values, points near 45° indicate good fit.

### 6.3 Learning Rate Comparison

Comparison of cost history for LR=1.0, 0.01, 0.0001.

## 7. Observations

- **Learning Curve:** Converges by ~1000 iterations.
- **Actual vs Predicted:**  $R^2 \approx 0.64$  means 64% variance explained.
- **Regularization:** Ridge ( $\lambda=10$ ) slightly improved test MSE and  $R^2$ .
- **Learning Rates:** 0.01 was optimal; 1.0 diverged, 0.0001 slow.

## **8. Conclusions**

Implementation meets assignment requirements. Regularization and hyperparameter experiments documented.

### **References:**

- [1] GeeksforGeeks on Gradient Descent
- [2] FutureAI blog on MSE and R<sup>2</sup>
- [3] GeeksforGeeks on MSE[3]
- [4] GfG on R<sup>2</sup>[4]